

How do Switches Maintain Routing Tables?

Distance-vector (Bellman-Ford)

Distance Vector (Bellman-Ford)
 $D_i^h \triangleq$ shortest distance to some destination
 s.t. path length $\leq h$ hops
 $G=(V,E)$
 $|V|=N$ nodes
 Init $D_{dest}^0 = 0$ $D_i^0 = \infty$ $i \neq dest$
 for $h=1,2,3,\dots,N-1$
 $D_i^h = \min_{k \in \text{Neighbors}(i)} \{D_k^{h-1} + w_{ik}\}$ $\forall i$
 $O(N^3)$

Matta @ BUCS - Routing 1-18

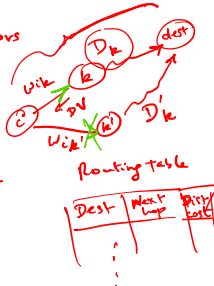
18

How do Switches Maintain Routing Tables?

Distance-vector (Bellman-Ford)

Periodically, send your neighbors
 Distance Vector $\rightarrow \{(dest, distance)\}$

- 1) unknown destination
 \Rightarrow take k as next hop
- 2) k is offering a shorter distance
 switch to k as next hop
- 3) k is my current next hop
 \Rightarrow update my own distance



Routing table

Dest	Next Hop	Cost
...

Matta @ BUCS - Routing 1-19

19

How do Switches Maintain Routing Tables?

Distance-vector (Bellman-Ford)

- Each node keeps in its routing table its current best distance (cost) to each destination node: (Destination, Cost, NextHop)
- Initially, the routing table of each node contains its neighbor nodes and distances to them
- Each node sends a copy of its routing information $\{(Destination, Cost)\}$ to each of its neighbor nodes
 - periodically (on the order of several seconds)
 - whenever its table changes (called **triggered** update)
- Upon receiving a message from a neighbor J , a node K updates its routing table as follows:
 - If node J knows a shorter path to some destination, node K updates its corresponding entry with J as the new next-hop
 - If node J reports a distance for an unknown destination, node K adds a new entry for it with J as the next-hop
 - If node K is using node J as its next-hop to some destination node and J reports a change in its distance to it, K updates the corresponding entry

Matta @ BUCS - Routing 1-20

20


Route Calculation in Distance Vector

```
void mergeRoute (Route *new) {
    int i;
    for (i = 0; i < numRoutes; ++i) {
        if (new->Dest == rt[i].Dest) {
            if (new->Cost + 1 < rt[i].Cost) // min-hop routing
                break;
            else if (new->NextHop == rt[i].NextHop)
                break;
            else
                return;
        }
    }
    rt[i] = *new;
    rt[i].TTL = MAX_TTL;
    ++rt[i].Cost;
    if (i == numRoutes)
        ++numRoutes;
}
```

Matta @ BUCS - Routing 1-21

21

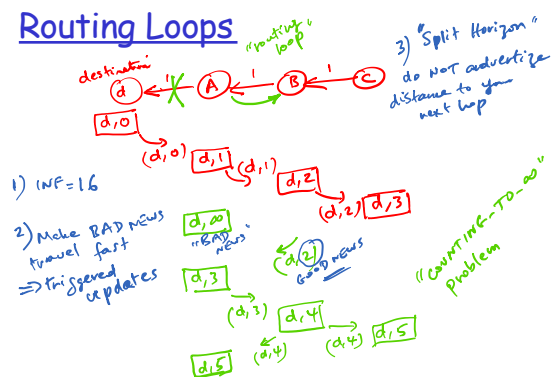
DV Routing



Matta @ BUCS - Routing 1-22

22

Routing Loops




Matta @ BUCS - Routing 1-23

23

Routing Loops

4) "Split Horizon with Poison Reverse"

No waiting loops involving 2 nodes




Matta @ BUCS - Routing 1-24

24

Routing Loops

PATH VECTOR

Advertise ∞
to your neighbor
if it's on the path
append node id
in routing message



Matta @ BUCS - Routing 1-25

25

Routing Loops

Figure 1 illustrates the execution of a distributed algorithm on a 5-processor system (A, B, C, D, E). The diagram shows the state of the processors at different stages of the algorithm, with a focus on processor A's interactions.

Initial State: All processors (A, B, C, D, E) are in the 'Initially' state.

Exchanges: The diagram shows a sequence of exchanges between processor A and other processors:

- After 1 exchange:** Processor A has exchanged with B.
- After 2 exchanges:** Processor A has exchanged with B and C.
- After 3 exchanges:** Processor A has exchanged with B, C, and D.
- After 4 exchanges:** Processor A has exchanged with B, C, D, and E.


The diagram also shows the state of the other processors (B, C, D, E) at each stage, indicating that they have also participated in exchanges. The final state shown is 'Initially', suggesting a reset or a specific configuration after the sequence of exchanges.

- Heuristics to break routing loops^(B)
 - m set infinity to maximum distance/cost
 - m split horizon: a node does not advertise to its next-hop
 - m split horizon with poison reverse: a node advertizes to its next-hop a distance of infinity

Matta @ BUCS - Routing 1-26

26

Routing Loops can still happen!




- link CD goes down
- loop forms involving A, B and C

Matta @ BUCS - Routing 1-27

27

Path Vector Routing



- loops can be completely avoided

Matta @ BUCS - Routing 1-28

28

Distance-Vector versus Link-State

Distance-Vector:

- Easy to implement
- Larger routing update messages: message size is proportional to the number of nodes in the network
- Slow to converge: route computation is distributed
- Loops/count-to-infinity may happen
- If link changes don't affect shortest path, no message exchange

Matta @ BUCS - Routing 1-29

29

Distance-Vector versus Link-State

Link-State:

- ❑ Smaller routing update messages: message size depends on the number of neighbors a node has
- ❑ Converges quickly: route computation is centralized
- ❑ A node stores a complete view of the network
- ❑ Any link change requires a broadcast

Both have strengths and weaknesses.

One or the other is used in almost every network

Matta @ BUCS - Routing 1-30
