

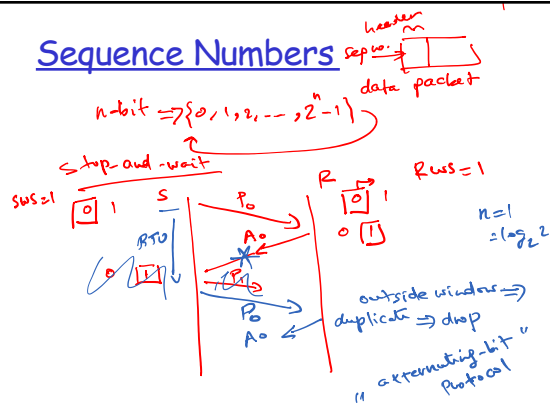
Sliding Window (cont'd)

- ❑ With Go-Back-N, **RWS** = 1
- ❑ With Selective Repeat, **RWS** = **SWS**.
Receiver can then maintain sequence numbers of packets that the sender can send, and so can detect whether a received packet is new or duplicate

Matta @ BUCS - Transport 1-36

36

Sequence Numbers




Matta @ BUCS - Transport 1-37

37

Sequence Numbers


- ❑ **SeqNum** field is finite; sequence numbers wrap around
- ❑ The size of the sequence number space must be larger than the number of outstanding packets
- ❑ Stop-and-Wait: sequence numbers $\{0, 1\}$



Matta @ BUCS - Transport 1-38

38

Sequence Numbers



Matta @ BUCS - Transport 1-39

39

Sequence Numbers (cont'd)

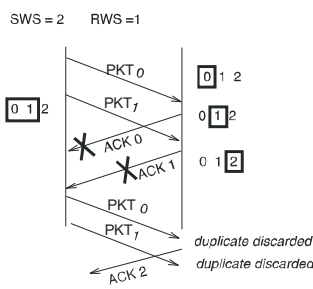
- Go-Back-N: sequence numbers $\{0, 1, \dots, SWS\}$

Matta @ BUCS - Transport 1-40

40

Sequence Numbers (cont'd)

- Go-Back-N: sequence numbers $\{0, 1, \dots, SWS\}$



Matta @ BUCS - Transport 1-41

41

Sequence Numbers (cont'd)


- Selective Repeat: sequence numbers $\{0, 1, \dots, SWS\}$ is not sufficient
 $SWS + RWS = 1$
 $SWS = 2$
 $RWS = 2$
 $\# \text{ bits for seq no.} = \lceil \log_2(SWS + RWS) \rceil$
-
- accept duplicate P_0
drop duplicates

Matta @ BUCS - Transport 1-42

42

Sequence Numbers (cont'd)

- Selective Repeat: sequence numbers $\{0, 1, \dots, SWS\}$ is not sufficient
 $SWS = 2$ $RWS = 2$
- Size of sequence number space must be at least $SWS + RWS = 2 SWS$
- Intuitively, SeqNum ``slides'' between two halves of sequence number space



Matta @ BUCS - Transport 1-43

43

End-to-End Challenges

Based basically on a reliable sliding window protocol, but it's challenging!

- Potentially different RTT
 - need adaptive timeout mechanism
- Potentially long delay in network
 - need to be prepared for arrival of very old packets
- Potentially different buffering at destination
 - need to accommodate different amounts of buffering
- Potentially different network capacity
 - need to be prepared for network congestion

Matta @ BUCS - Transport 1-44

44

TCP Segment Format

- Demultiplexing: each connection identified with 4-tuple:


- $\langle \text{SrcPort}, \text{SrcIPAddr}, \text{DstPort}, \text{DstIPAddr} \rangle$
- IP addresses obtained from IP layer!

- Basically, a sliding window operating at byte (not segment) level

- Acknowledgment, SequenceNum, AdvertisedWindow
- Piggybacking ACK on data segments destined for sender improves network utilization

- Flags

- SYN, FIN, RESET, PUSH, URG, ACK




Matta @ BUCS - Transport 1-45

45

TCP Reliability & Flow Control

- Like SR with Explicit Rtx/cumulative ACKs:

- storing *out-of-order* bytes
- using *one timer* for all unacked bytes
- using *duplicate ACK* to fast retransmit
- On retransmission, *only one segment retransmitted*



Matta @ BUCS - Transport 1-46

46

TCP Reliability & Flow Control

- Like SR with Explicit Rtx/cumulative ACKs:

- storing *out-of-order* bytes
- using *one timer* for all unacked bytes
- using *duplicate ACK* to fast retransmit
- On retransmission, *only one segment retransmitted*

- A new version, with SACK option, is more like GBN with selective repeats!

- At sender:

- $\text{LastByteSent} - \text{LastByteAcked} \leq \text{AdvertizedWindow}$
- If zero, sender keeps sending 1-byte data segments

Matta @ BUCS - Transport 1-47

47