

22

---

---

---

---

---

---

---

---

### Socket Operations in C/C++

- Creating a socket
 

```
int socket(int domain, int type, int protocol)
```

  - domain=AF\_INET (for TCP/IP protocols)
  - type=SOCK\_STREAM (for TCP-based application)
- Passive open on server
 

```
int bind(int socket, struct sockaddr *address, int addr_len)
int listen(int socket, int backlog)
int accept(int socket, struct sockaddr *address, int addr_len)
```

*client's*

*IP address (4B)  
port number (2B)*

Matta © BUCS - Applications 1-23

23

---

---

---

---

---

---

---

---

### Socket Operations (cont'd)

- Active open on client
 

```
int connect(int socket, struct sockaddr *address, int addr_len)
```

*server's*
- Sending and receiving messages
 

```
int send(int socket, char *message, int msg_len, int flags)
int recv(int socket, char *buffer, int buf_len, int flags)
```

Matta © BUCS - Applications 1-24

24

---

---

---

---

---

---

---

---

## Concurrent Server Implementation

Main server does the following:

- ❑ Open port:
  - ▢ Main server opens well-known port
- ❑ Wait for client:
  - ▢ Main server waits for a new client request
- ❑ Start Worker:
  - ▢ Main server (parent) starts a worker (child) to handle request (e.g., in UNIX/Linux, it **forks** a copy of the server process)
  - ▢ parent closes ``connected`` socket, and child closes ``listening`` socket
  - ▢ child dies when done
- ❑ Continue:
  - ▢ parent returns to the **wait** step

Matta @ BUCS - Applications 1-25

25

## Server Code: establish socket

```
/* code to establish a socket */
int establish(unsigned short portnum) {
    char myname[MAXHOSTNAME+1];
    int s;
    struct sockaddr_in sa;
    struct hostent *hp;

    memset(&sa, 0, sizeof(struct sockaddr)); /* clear our address */
    gethostname(myname, MAXHOSTNAME); /* who are we? */
    hp = gethostbyname(myname); /* get our address info */
    if (hp == NULL) /* we don't exist! */
        return(-1);
    sa.sin_family = hp->h_addrtype; /* this is our host address */
    sa.sin_addr = htonl(INADDR_ANY); /* this is our default IP address */
    sa.sin_port = htons(portnum); /* this is our port number */
    if ((s = socket(AF_INET, SOCK_STREAM, 0)) < 0) /* create socket */
        return(-1);
    if (bind(s, (struct sockaddr *)&sa, sizeof(sa)) < 0) {
        close(s);
        return(-1); /* bind address to socket */
    }
    listen(s, 3); /* max # of queued connects */
    return(s);
}
```

Matta @ BUCS - Applications 1-26

26

## Server Code: wait for clients

```
/* wait for a connection to occur on a socket created with establish() */
int get_connection(int s) {
    int t; /* socket of connection */

    if ((t = accept(s, NULL, NULL)) < 0) /* accept connection if there is one */
        return(-1);
    return(t);
}
```

Matta @ BUCS - Applications 1-27

27

## Server Code: main program

```
#include <errno.h> /* obligatory includes */
#include <signal.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <sys/wait.h>
#include <netinet/in.h>
#include <netdb.h>

#define PORTNUM 5000 /* random port number, we need something */
void do_something(int);

main() {
    int s, t;

    if ((s = establish(PORTNUM)) < 0) { /* plug in the phone */
        perror("establish");
        exit(1);
    }
}
```

Matta @ BUCS - Applications 1-28

28

## Server Code: main (cont' d)

```
for (;;) { /* loop for phone calls */
    if ((t = get_connection(s)) < 0) { /* get a connection */
        perror("accept"); /* bad */
        exit(1);
    }
    do_something(t);
    close(t); /* another connection */
    continue;
} /* end of main */

/* this is the function that plays with the socket. It will
   be called after getting a connection. */
void do_something(int t) {
    /* do your thing with the socket here */
}
```

Matta @ BUCS - Applications 1-29

29

## Concurrent Server

```
/* how a concurrent server looks like */
for (;;) { /* loop for phone calls */
    if ((t = get_connection(s)) < 0) { /* get a connection */
        perror("accept"); /* bad */
        exit(1);
    }
    switch( fork() ) { /* try to handle connection */
        case -1: /* bad news. scream and die */
            perror("fork");
            close(s);
            close(t);
            exit(1);
        case 0: /* we're the child, do something */
            close(s);
            do_something(t);
            close(t);
            exit(0);
        default: /* we're the parent so look for */
            /* another connection */
            continue;
    }
}
```

Matta @ BUCS - Applications 1-30

30

## Client Code

```
int call_socket(char *hostname, unsigned short portnum) {
    struct sockaddr_in sa;
    struct hostent *hp;
    int a, c;

    if ((hp = gethostbyname(hostname)) == NULL) { /* do we know the host's address? */
        return(-1); /* no */
    }
    memset(&sa, 0, sizeof(sa));
    memcpy((char *)&sa.sin_addr, hp->h_addr, hp->h_length); /* set address */
    sa.sin_family = hp->h_addrtype;
    sa.sin_port = htons(portnum);
    if ((c = socket(hp->h_addrtype, SOCK_STREAM, 0)) < 0) /* get socket */
        return(-1);
    if (connect(c, (struct sockaddr *)&sa, sizeof(sa)) < 0) { /* connect */
        close(s);
        return(-1);
    }
    return(c);
}
```

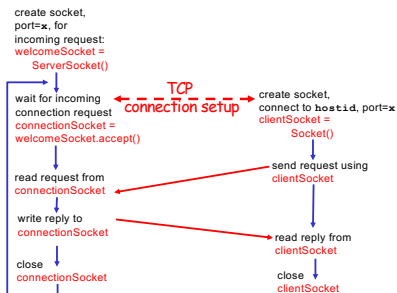
Matta @ BUCS - Applications 1-31

31

## Java client/server socket interaction: TCP

Server (running on hostid)

Client



Matta @ BUCS - Applications 1-32

32

## Example: Java client (TCP)

```
import java.io.*;
import java.net.*;
class TCPClient {

    public static void main(String argv[]) throws Exception
    {
        String sentence;
        String modifiedSentence;

        Create
        input stream → BufferedReader inFromUser =
                       new BufferedReader(new InputStreamReader(System.in));
        Create
        client socket, → Socket clientSocket = new Socket("hostname", 6789);
        connect to server
        Create
        output stream → DataOutputStream outToServer =
                       new DataOutputStream(clientSocket.getOutputStream());
        attached to socket
    }
}
```

Matta @ BUCS - Applications 1-33

33

### Example: Java client (TCP), cont.

```

    Create
    input stream
    attached to socket } BufferedReader inFromServer =
                        new BufferedReader(new
                        InputStreamReader(clientSocket.getInputStream()));

    Send line
    to server } sentence = inFromUser.readLine();

    Read line
    from server } outToServer.writeBytes(sentence + '\n');
                modifiedSentence = inFromServer.readLine();
                System.out.println("FROM SERVER: " + modifiedSentence);
                clientSocket.close();
            }
        }
    }
}

```

Matta @ BUCS - Applications 1-34

34

### Example: Java server (TCP)

```
import java.io.*;
import java.net.*;

class TCPServer {

    public static void main(String argv[]) throws Exception
    {
        String clientSentence;
        String capitalizedSentence;

        ServerSocket welcomeSocket = new ServerSocket(6789);

        while(true) {

            Socket connectionSocket = welcomeSocket.accept();

            BufferedReader inFromClient =
                new BufferedReader(new
                    InputStreamReader(connectionSocket.getInputStream()));
        }
    }
}
```

Matta @ BUCS - Applications 1-35

35

### Example: Java server (TCP), cont

```

    DataOutputStream outToClient =
        new DataOutputStream(connectionSocket.getOutputStream());

    clientSentence = inFromClient.readLine();

    capitalizedSentence = clientSentence.toUpperCase() + "\n";

    outToClient.writeBytes(capitalizedSentence);
    connectionSocket.close();
}
}

```

Create output stream, attached to socket  
 Read in line from socket  
 Write out line to socket  
 End of while loop, loop back and wait for another client connection

Matta @ BUCS - Applications 1-36

36

## Multi-threaded Programs

- ❑ A thread is a lightweight process
- ❑ A process can have one or more threads
- ❑ A thread runs in the context of a process
  - All threads share access to code and data, but each thread has its own private PC, registers, stack and state
- ❑ A server would have a thread to handle each request
- ❑ A client could also have multiple threads, e.g., one to send requests to server and another to receive responses from server
- ❑ Java threads (discussed in lab)

Matte @ BUCS - Applications 1-37

37

## Client/server Socket Interaction in C/C++: UDP

Server (running on *hostid*)

```
create socket,
bind it to port*, for
incoming request:
socket()
bind()
↓
read request
recvfrom()
↓
write reply
sendto()
specifying client
host address,
port number
```

Client

```
create socket, bind it
socket()
bind()
↓
create address (hostid, port=x),
send datagram request
using sendto()
↓
read reply
recvfrom()
↓
close socket
close()
```

Matte @ BUCS - Applications 1-38

38


## Caching example:

### assumptions:

- ❖ avg object size: 100,000 bits
- ❖ avg request rate from browsers to origin servers: 15/sec
- ❖ avg data rate to browsers: 1.50 Mbps
- ❖ RTT over the Internet: 2 sec
- ❖ access link rate: 1.51 Mbps

### consequences:

- ❖ LAN utilization: 0.15%
  - ❖ Assume LAN delay ~ usecs
- ❖ access link utilization ≈ 99% *problem!*
- ❖ total delay = Internet delay + access delay + LAN delay = 2 sec + *minutes* + *usecs*



39

## Caching example: fatter access link

### assumptions:

- ❖ avg object size: 100K bits
- ❖ avg request rate from browsers to origin servers: 15/sec
  - ❖ avg data rate to browsers: 1.50 Mbps
- ❖ RTT over the Internet: 2 sec


- ❖ access link rate: 1.51 Mbps

### consequences:

- ❖ LAN utilization: 0.15%
- ❖ access link utilization = 99%
- ❖ total delay = Internet delay + access delay + LAN delay

$$= 2 \text{ sec} + \text{minutes} + \text{usecs} \approx 2 \text{ sec}$$

Cost: increased access link speed (not cheap!)



Application Layer 2-40

40

## Caching example: install local cache

### assumptions:

- ❖ avg object size: 100K bits
- ❖ avg request rate from browsers to origin servers: 15/sec
  - ❖ avg data rate to browsers: 1.50 Mbps
- ❖ RTT over the Internet: 2 sec


- ❖ access link rate: 1.51 Mbps

### consequences:

- ❖ LAN utilization: 0.15%
- ❖ access link utilization = ?
- ❖ total delay = ?

How to compute link utilization, delay?

Cost: web cache (cheap!)




Application Layer 2-41

41

## Caching example: install local cache

### Calculating access link utilization, delay with cache:

- ❑ suppose cache hit rate is 0.4
  - 40% requests satisfied at cache,
  - 60% requests satisfied at origin




Application Layer 2-42

42

## Caching example: install local cache

### Calculating access link utilization, delay with cache:


- suppose cache hit rate is 0.4
  - 40% requests satisfied at cache, 60% requests satisfied at origin
- ❖ access link utilization:
  - 60% of requests use access link
- ❖ data rate to browsers over access link =  $0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$ 
  - utilization =  $0.9 / 1.51 = .6$
  - Assume access delay ~ 700ms
- ❖ total delay
  - =  $0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
  - =  $0.6 * (\sim 2.7) + 0.4 * (\sim \text{usecs})$
  - =  $\sim 1.6 \text{ secs}$
  - less than with 151 Mbps link (and cheaper too!)



Application Layer 2-43

43

## Application Example: TCP/IP Email



Matta @ BUCS - Applications 1-44

44